
MODULE *EHS_ChaosPrevention*

EHS_ChaosPrevention is a spec for automatic failover protection mechanism in *Ehs EHS* (*EdgeHealthService*) is a monitoring service running in each of Azure *Frontdoor*'s Edge which determines whether that Edge can serve user traffic or not. When an Edge is determined not fit to serve traffic then it is turned off. In order to prevent global outage of all *Edges*, we need this protection mechanism to prevent all *EHS* instances from turning off all *Edges* at the same time.

EXTENDS *FiniteSets*, *Integers*

CONSTANT

Edges, Set of *Edges* serving user traffic

NumEdgesAllowedToTurnOff Number of *Edges* allowed to turn off automatically.

ASSUME $\text{NumEdgesAllowedToTurnOff} \leq \text{Cardinality}(\text{Edges})$

VARIABLES

edgeState, State of *Edges* indicating whether it can serve traffic or not.

ehsDecision, *EHS* decision for the respective edge

messagingState, Messaging state of an edge indicating whether it has sent its message to other edges or not.

msgs Messages sent to other *Edges*

State Filters

$\text{EdgesInOffState} \triangleq \{e \in \text{Edges} : \text{edgeState}[e] = \text{"Off"}\}$

$\text{EdgesInTurningOffState} \triangleq \{e \in \text{Edges} : \text{edgeState}[e] = \text{"TurningOff"}\}$

Model Messages

$\text{Messages} \triangleq$

$[type : \{\text{"StatusRequest"}\}, sender : \text{EdgesInTurningOffState}, receiver : \text{Edges}]$

\cup

$[type : \{\text{"StatusResponse"}\}, sender : \text{Edges}, receiver : \text{EdgesInTurningOffState}, val : \{\text{"On"}, \text{"TurningOff"}, \text{"Off"}\}]$

Model Init

$\text{EhsInit} \triangleq$

$\wedge \text{edgeState} = [edge \in \text{Edges} \mapsto \text{"On"}]$

$\wedge \text{ehsDecision} = [ehs \in \text{Edges} \mapsto \text{"Healthy"}]$

$\wedge \text{messagingState} = [ehs \in \text{Edges} \mapsto \text{"Reset"}]$

$\wedge \text{msgs} = \{\}$

Message Actions

$\text{SendMessage}(m) \triangleq \text{msgs}' = \text{msgs} \cup \{m\}$

$\text{SendMessages}(m) \triangleq \text{msgs}' = \text{msgs} \cup m$

Message Filtering

$HaveGotAllStatusResponseMessages(e) \triangleq$

Determines whether an Edge has got responses from all the *Edges* to which it sent the request to.

LET $allSentMessages \triangleq \{m \in msgs : \wedge m.type = \text{"StatusRequest"} \wedge m.sender = e\}$

$receivedMessagesForLatestSeq \triangleq \{m \in msgs : \wedge m.type = \text{"StatusResponse"} \wedge m.receiver = e\}$

IN $\wedge Cardinality(allSentMessages) = Cardinality(receivedMessagesForLatestSeq)$

$CompleteStatusRequestMessage(e) \triangleq$

Complete status request message seq once by removing the request & response messages for a sequence

LET $messagesBelongingToEdge \triangleq \{m \in msgs : \vee (m.type = \text{"StatusRequest"} \wedge m.sender = e) \vee (m.type = \text{"StatusResponse"} \wedge m.receiver = e)\}$

IN $\wedge msgs' = msgs \setminus messagesBelongingToEdge$

$EdgesThatRespondWithGivenStateForLatestRequest(state, e) \triangleq$

Returns that responded with the given state (*On*, *Off*, *Turning off* or *Timeout*) for the latest Status Request message

$\{m1.sender : m1 \in \{m \in msgs : \wedge m.type = \text{"StatusResponse"} \wedge m.receiver = e \wedge m.val = state\}\}$

$StateOfTargetEdgeAsPerEdge(targetEdge, e) \triangleq$

Returns what is the state of the edge 'TargetEdge' in the view of edge 'e' based on the response of 'TargetEdge' to the latest *StatusRequest* message from edge 'e'

$\{m1.val : m1 \in \{m \in msgs : \wedge m.type = \text{"StatusResponse"} \wedge m.receiver = e \wedge m.sender = targetEdge\}\}$

$EdgesNotInOnStateInTheViewOfEdge(e) \triangleq$

Get the list of *Edges* which are not in the *On* state. That is any edge which responded with *Off*, *Turning Off* or *Timeout* as status. Also, include the edge itself.

$EdgesThatRespondWithGivenStateForLatestRequest(\text{"Off"}, e)$

\cup

$EdgesThatRespondWithGivenStateForLatestRequest(\text{"TurningOff"}, e)$

\cup

$EdgesThatRespondWithGivenStateForLatestRequest(\text{"Timeout"}, e)$

\cup

IF $edgeState[e] \neq \text{"On"}$ THEN $\{e\}$

ELSE $\{\}$

$CanTurnOff(e) \triangleq$

An Edge can turn off only if (total number of edges which responded as *Off*, *TurnedOff* or *Timeout*) \leq Number *Edges* Allowed to Turn *Off*

$\wedge Cardinality(EdgesNotInOnStateInTheViewOfEdge(e)) \leq NumEdgesAllowedToTurnOff$

State Changes

$EhsHealthyToUnHealthy(e) \triangleq$

EHS can at anytime determine that edge goes from *Healthy* to *UnHealthy* state
 $\wedge ehsDecision[e] = \text{"Healthy"}$
 $\wedge ehsDecision' = [ehsDecision \text{ EXCEPT } ![e] = \text{"Unhealthy"}]$
 $\wedge \text{UNCHANGED } \langle edgeState, msgs, messagingState \rangle$

$EhsUnHealthyToHealthy(e) \triangleq$

EHS can at anytime determine that edge goes from *UnHealthy* to *Healthy* state
 $\wedge ehsDecision[e] = \text{"Unhealthy"}$
 $\wedge ehsDecision' = [ehsDecision \text{ EXCEPT } ![e] = \text{"Healthy"}]$
 $\wedge \text{UNCHANGED } \langle edgeState, msgs, messagingState \rangle$

$EdgeOnToTurningOff(e) \triangleq$

When Ehs' decision is *Unhealthy* then the Edge goes to *TurningOff* state if it is in *On* state.
 $\wedge ehsDecision[e] = \text{"Unhealthy"}$
 $\wedge edgeState[e] = \text{"On"}$
 $\wedge edgeState' = [edgeState \text{ EXCEPT } ![e] = \text{"TurningOff"}]$
 $\wedge \text{UNCHANGED } \langle ehsDecision, msgs, messagingState \rangle$

$EdgeTurningOffToOn(e) \triangleq$

When Ehs' decision is *Healthy* then the Edge goes to *On* state from *Turning Off* state, provided it has completed its messaging process which it started when it went to *Turning Off* state
 $\wedge ehsDecision[e] = \text{"Healthy"}$
 $\wedge edgeState[e] = \text{"TurningOff"}$
 $\wedge messagingState[e] = \text{"Completed"}$
 $\wedge edgeState' = [edgeState \text{ EXCEPT } ![e] = \text{"On"}]$
 $\wedge messagingState' = [messagingState \text{ EXCEPT } ![e] = \text{"Reset"}]$
 $\wedge CompleteStatusRequestMessage(e)$
 $\wedge \text{UNCHANGED } \langle ehsDecision \rangle$

$EdgeTurningOffToTurningOff(e) \triangleq$

This is a case where the Edge can't move away from the *Turning Off* state because there are too many edges in turning off or off state or have not responded to the status request message
 $\wedge ehsDecision[e] = \text{"Unhealthy"}$
 $\wedge edgeState[e] = \text{"TurningOff"}$
 $\wedge messagingState[e] = \text{"Completed"}$
 $\wedge \neg CanTurnOff(e)$
 $\wedge messagingState' = [messagingState \text{ EXCEPT } ![e] = \text{"Reset"}]$
 $\wedge CompleteStatusRequestMessage(e)$
 $\wedge \text{UNCHANGED } \langle edgeState, ehsDecision \rangle$

$EdgeTurningOffToOff(e) \triangleq$

Edge has gotten responses from other *Edges* and it determines that the total number of edges not in *On* state are below the allowed threshold. Hence, it can go to *Off* State from *Turning Off* state as it is still unhealthy.
 $\wedge ehsDecision[e] = \text{"Unhealthy"}$

$\wedge \text{edgeState}[e] = \text{"TurningOff"}$
 $\wedge \text{messagingState}[e] = \text{"Completed"}$
 $\wedge \text{CanTurnOff}(e)$
 $\wedge \text{edgeState}' = [\text{edgeState} \text{ EXCEPT } ![e] = \text{"Off"}]$
 $\wedge \text{messagingState}' = [\text{messagingState} \text{ EXCEPT } ![e] = \text{"Reset"}]$
 $\wedge \text{CompleteStatusRequestMessage}(e)$
 $\wedge \text{UNCHANGED } \langle \text{ehsDecision} \rangle$

$\text{EdgeOffToOn}(e) \triangleq$

When Ehs' decision is *Healthy* then the Edge goes to *On* state if it is in *Off* state.

$\wedge \text{ehsDecision}[e] = \text{"Healthy"}$
 $\wedge \text{edgeState}[e] = \text{"Off"}$
 $\wedge \text{edgeState}' = [\text{edgeState} \text{ EXCEPT } ![e] = \text{"On"}]$
 $\wedge \text{UNCHANGED } \langle \text{ehsDecision}, \text{msgs}, \text{messagingState} \rangle$

Messaging Actions

$\text{NotRespondedToStatusRequestMessage}(e, \text{statusRequestMessage}) \triangleq$

$\wedge \neg(\exists m \in \text{msgs} : \wedge m.\text{type} = \text{"StatusResponse"}$
 $\wedge m.\text{sender} = e$
 $\wedge m.\text{receiver} = \text{statusRequestMessage.sender})$

$\text{EdgeProcessResponses}(e) \triangleq$

Once an Edge has received all the responses for its requests then the message has completed.

$\wedge \text{edgeState}[e] = \text{"TurningOff"}$
 $\wedge \text{ehsDecision}[e] = \text{"Unhealthy"}$
 $\wedge \text{messagingState}[e] = \text{"Sent"}$
 $\wedge \text{HaveGotAllStatusResponseMessages}(e)$
 $\wedge \text{messagingState}' = [\text{messagingState} \text{ EXCEPT } ![e] = \text{"Completed"}]$
 $\wedge \text{UNCHANGED } \langle \text{ehsDecision}, \text{edgeState}, \text{msgs} \rangle$

$\text{EdgeSendsRequest}(e) \triangleq$

When an edge comes to *Turning Off*, then it will send a request to all other edges asking their state.

$\wedge \text{edgeState}[e] = \text{"TurningOff"}$
 $\wedge \text{ehsDecision}[e] = \text{"Unhealthy"}$
 $\wedge \text{messagingState}[e] = \text{"Reset"}$
 $\wedge \text{messagingState}' = [\text{messagingState} \text{ EXCEPT } ![e] = \text{"Sent"}]$
 $\wedge \text{SendMessage}(\{[type \mapsto \text{"StatusRequest"}, \text{sender} \mapsto e, \text{receiver} \mapsto e1] : e1 \in \text{Edges} \setminus \{e\}\})$
 $\wedge \text{UNCHANGED } \langle \text{ehsDecision}, \text{edgeState} \rangle$

$\text{EdgeRespondsWithStatus}(e) \triangleq$

Whenever there is a pending request, asking its state, for the edge then the edge responds with its state.

$\wedge \exists m \in \text{msgs} :$
 $\wedge m.\text{type} = \text{"StatusRequest"}$
 $\wedge m.\text{receiver} = e$
 $\wedge \text{NotRespondedToStatusRequestMessage}(e, m)$

$\wedge \text{SendMessage}([type \mapsto \text{"StatusResponse"}, sender \mapsto e, receiver \mapsto m.sender, val \mapsto edgeState[e]])$
 $\wedge \text{UNCHANGED } \langle ehsDecision, edgeState, messagingState \rangle$

$\text{EdgeRespondsWithAbort}(e) \triangleq$

At times an edge can respond with a timeout message for pending state request to it.

This simulates request timeouts

$\wedge \exists m \in msgs :$
 $\wedge m.type = \text{"StatusRequest"}$
 $\wedge m.receiver = e$
 $\wedge \text{NotRespondedToStatusRequestMessage}(e, m)$
 $\wedge \text{SendMessage}([type \mapsto \text{"StatusResponse"}, sender \mapsto e, receiver \mapsto m.sender, val \mapsto \text{"Timeout"}])$
 $\wedge \text{UNCHANGED } \langle ehsDecision, edgeState, messagingState \rangle$

Model Invariants

$\text{EhsTypeOK} \triangleq$

Invariant ensuring that the different state don't take an invalid value.

$\wedge edgeState \in [Edges \rightarrow \{\text{"On"}, \text{"TurningOff"}, \text{"Off"}\}]$
 $\wedge ehsDecision \in [Edges \rightarrow \{\text{"Healthy"}, \text{"Unhealthy"}\}]$
 $\wedge messagingState \in [Edges \rightarrow \{\text{"Reset"}, \text{"Sent"}, \text{"Completed"}\}]$

$\text{EhsStateOK} \triangleq$

Important Invariant which ensures that not too many *Edges* are in *Off* state thus preventing global outage.

$\wedge \text{Cardinality}(EdgesInOffState) \leq \text{NumEdgesAllowedToTurnOff}$

$\text{EhsModelOk} \triangleq$

$\wedge \text{EhsTypeOK}$
 $\wedge \text{EhsStateOK}$

$\text{NoEdgeHasLiedAboutItsState} \triangleq$

Invariant which ensures that no edges gives a wrong reply to the *StatusRequest* Message.

$\neg(\exists m \in msgs : \wedge m.type = \text{"StateResponse"}$
 $\wedge (m.val \neq \text{"Timeout"} \wedge edgeState[m.sender] \neq m.val))$

$\text{NoTwoTurningOffEdgesSeeEachOtherInOnState} \triangleq$

Invariant which ensures there are no race condition when two edges are turning off at the same time.

No two edges which are going to turning off state at the time should see each other in *On* state.

This ensures that the edges updates its state before requesting other edges about their respective state.

$\neg(\exists e1, e2 \in Edges : \wedge e1 \neq e2$
 $\wedge edgeState[e1] = \text{"TurningOff"}$
 $\wedge edgeState[e2] = \text{"TurningOff"}$
 $\wedge messagingState[e1] = \text{"Completed"}$
 $\wedge messagingState[e2] = \text{"Completed"}$
 $\wedge \text{StateOfTargetEdgeAsPerEdge}(e2, e1) = \{\text{"On"}\}$
 $\wedge \text{StateOfTargetEdgeAsPerEdge}(e1, e2) = \{\text{"On"}\})$

$$\begin{aligned}
EhsNext &\triangleq \\
&\vee \exists e \in Edges : \\
&\quad EhsHealthyToUnHealthy(e) \vee EhsUnHealthyToHealthy(e) \\
&\quad \vee EdgeOnToTurningOff(e) \vee EdgeTurningOffToOn(e) \vee EdgeTurningOffToOff(e) \vee EdgeTurningOff \\
&\quad \vee EdgeSendsRequest(e) \vee EdgeProcessResponses(e) \vee EdgeRespondsWithStatus(e) \vee EdgeRespond \\
EhsSpec &\triangleq EhsInit \wedge \Box[EhsNext]_{\langle edgeState, ehsDecision, messagingState, msgs \rangle}
\end{aligned}$$

\ * Modification History
\ * Last modified Wed May 04 19:02:51 PDT 2016 by guhanr
\ * Created Wed May 04 19:02:11 PDT 2016 by guhanr