
MODULE *TwoPhase*

This specification describes the Two-Phase *Commit* protocol, in which a transaction manager (*TM*) coordinates the resource managers (*RM*s) to implement the Transaction *Commit* specification of module *TCommit*. In this specification, *RM*s spontaneously issue *Prepared* messages. We ignore the *Prepare* messages that the *TM* can send to the *RM*s.

For simplicity, we also eliminate *Abort* messages sent by an *RM* when it decides to abort. Such a message would cause the *TM* to abort the transaction, an event represented here by the *TM* spontaneously deciding to abort.

This specification describes only the safety properties of the protocol—that is, what is allowed to happen. What must happen would be described by liveness properties, which we do not specify.

CONSTANT *RM* The set of resource managers

VARIABLES

<i>rmState</i> ,	<i>rmState</i> [<i>r</i>] is the state of resource manager <i>r</i> .
<i>tmState</i> ,	The state of the transaction manager.
<i>tmPrepared</i> ,	The set of <i>RM</i> s from which the <i>TM</i> has received “Prepared”
	messages.

msgs

In the protocol, processes communicate with one another by sending messages. Since we are specifying only safety, a process is not required to receive a message, so there is no need to model message loss. (There’s no difference between a process not being able to receive a message because the message was lost and a process simply ignoring the message.) We therefore represent message passing with a variable *msgs* whose value is the set of all messages that have been sent. Messages are never removed from *msgs*. An action that, in an implementation, would be enabled by the receipt of a certain message is here enabled by the existence of that message in *msgs*. (Receipt of the same message twice is therefore allowed; but in this particular protocol, receiving a message for the second time has no effect.)

Messages \triangleq

The set of all possible messages. Messages of type “Prepared” are sent from the *RM* indicated by the message’s *rm* field to the *TM*. Messages of type “Commit” and “Abort” are broadcast by the *TM*, to be received by all *RM*s. The set *msgs* contains just a single copy of such a message.

$$[type : \{ \text{“Prepared”} \}, rm : RM] \cup [type : \{ \text{“Commit”}, \text{“Abort”} \}]$$

TPTypeOK \triangleq

The type-correctness invariant

$$\begin{aligned} &\wedge rmState \in [RM \rightarrow \{ \text{“working”}, \text{“prepared”}, \text{“committed”}, \text{“aborted”} \}] \\ &\wedge tmState \in \{ \text{“init”}, \text{“committed”}, \text{“aborted”} \} \\ &\wedge tmPrepared \subseteq RM \\ &\wedge msgs \subseteq Messages \end{aligned}$$

TPInit \triangleq

The initial predicate.

$$\begin{aligned} &\wedge rmState = [r \in RM \mapsto \text{“working”}] \\ &\wedge tmState = \text{“init”} \\ &\wedge tmPrepared = \{ \} \\ &\wedge msgs = \{ \} \end{aligned}$$

We now define the actions that may be performed by the processes, first the *TM*'s actions, then the *RM*s' actions.

$TM_{RcvPrepared}(r) \triangleq$

The *TM* receives a "Prepared" message from resource manager *r*.

$\wedge tmState = \text{"init"}$
 $\wedge [type \mapsto \text{"Prepared"}, rm \mapsto r] \in msgs$
 $\wedge tmPrepared' = tmPrepared \cup \{r\}$
 $\wedge \text{UNCHANGED } \langle rmState, tmState, msgs \rangle$

$TM_{Commit} \triangleq$

The *TM* commits the transaction; enabled iff the *TM* is in its initial state and every *RM* has sent a "Prepared" message.

$\wedge tmState = \text{"init"}$
 $\wedge tmPrepared = RM$
 $\wedge tmState' = \text{"committed"}$
 $\wedge msgs' = msgs \cup \{[type \mapsto \text{"Commit"}]\}$
 $\wedge \text{UNCHANGED } \langle rmState, tmPrepared \rangle$

$TM_{Abort} \triangleq$

The *TM* spontaneously aborts the transaction.

$\wedge tmState = \text{"init"}$
 $\wedge tmState' = \text{"aborted"}$
 $\wedge msgs' = msgs \cup \{[type \mapsto \text{"Abort"}]\}$
 $\wedge \text{UNCHANGED } \langle rmState, tmPrepared \rangle$

$RM_{Prepare}(r) \triangleq$

Resource manager *r* prepares.

$\wedge rmState[r] = \text{"working"}$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{"prepared"}]$
 $\wedge msgs' = msgs \cup \{[type \mapsto \text{"Prepared"}, rm \mapsto r]\}$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared \rangle$

$RM_{ChooseToAbort}(r) \triangleq$

Resource manager *r* spontaneously decides to abort. As noted above, *r* does not send any message in our simplified spec.

$\wedge rmState[r] = \text{"working"}$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{"aborted"}]$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared, msgs \rangle$

$RM_{RcvCommitMsg}(r) \triangleq$

Resource manager *r* is told by the *TM* to commit.

$\wedge [type \mapsto \text{"Commit"}] \in msgs$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{"committed"}]$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared, msgs \rangle$

$RM_{RcvAbortMsg}(r) \triangleq$

Resource manager r is told by the TM to abort.

$\wedge [type \mapsto \text{"Abort"}] \in msgs$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{"aborted"}]$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared, msgs \rangle$

$TPNext \triangleq$
 $\vee TMCommit \vee TMAbort$
 $\vee \exists r \in RM :$
 $TMRcvPrepared(r) \vee RMPPrepare(r) \vee RMChooseToAbort(r)$
 $\vee RMRcvCommitMsg(r) \vee RMRcvAbortMsg(r)$

$TPSpec \triangleq TPInit \wedge \Box [TPNext]_{\langle rmState, tmState, tmPrepared, msgs \rangle}$

The complete spec of the Two-Phase *Commit* protocol.

THEOREM $TPSpec \Rightarrow \Box TPTypeOK$

This theorem asserts that the type-correctness predicate $TPTypeOK$ is an invariant of the specification.

We now assert that the Two-Phase *Commit* protocol implements the Transaction *Commit* protocol of module $TCommit$. The following statement defines $TC!TCSpec$ to be formula $TCSpec$ of module $TCommit$. (The TLA+ `INSTANCE` statement imports all the definitions from module $TCommit$ renamed in this way, thus avoiding any name conflicts that might exist between defined operators in the two modules. The constant RM and variable $rmState$ are the same in both modules.)

$TC \triangleq \text{INSTANCE } TCommit$

THEOREM $TPSpec \Rightarrow TC!TCSpec$

This theorem asserts that the specification $TPSpec$ of the Two-Phase Commit protocol implements the specification $TCSpec$ of the Transaction *Commit* protocol.

The two theorems in this module have been checked with *TLC* for six *RM*s, a configuration with 50816 reachable states, in a little over a minute on a 1 *GHz PC*.